



FreCoS: A Learned Staleness Model for Semantic Caches

A freshness- and cost-aware extension to GPTCache

Lior Lotan

Department of Software and Information Systems Engineering
Ben-Gurion University of the Negev

August 1, 2026

Abstract

Semantic caches are usually evaluated on hit rate, latency, and at best false-hit rate, none of which asks whether a hit was still true at the moment it was served. I name that gap *stale-hit-rate*, the fraction of served hits whose content had already expired, and extend GPTCache with a per-cluster staleness model used both as a validity gate on the serve path and as a decay term in eviction. This version replaces an earlier oracle-cluster, exact-match harness with real embedder-based clustering and a real semantic index, and that change reverses the report’s central finding: under real clustering, the adjusted Rand index between the k-means-on-embeddings assignment and the workload’s true clusters is 0.02–0.06, barely above chance, because the workload’s canonical query text differs across clusters only in two embedded numbers, and a general-purpose embedder cannot recover cluster identity from that. Per-cluster staleness fitting was therefore never actually exercised; “learned tracks global, not oracle” is a diagnosed measurement failure, not a verdict on per-cluster learning in general. What still holds: the validity gate cuts stale-hit-rate roughly tenfold regardless of clustering quality, since it applies the same (possibly wrong) cluster consistently at calibration and serve time; and a fair test of cost-aware eviction, run for the first time under conditions where eviction actually occurs, shows FreCoS beating a frequency floor on cost saved with a large, significant effect, though plain LRU beats FreCoS on both cost saved and stale-hit-rate on this trace. The same real semantic index also exposes a false-hit problem the old harness could not see: a real 0.8 cosine threshold on this workload’s near-identical templates serves the wrong answer roughly 97% of the time a hit is served at all: a limitation of this synthetic workload’s text, not of the cache mechanism.

1 Introduction

1.1 Motivation

A semantic cache sits in front of an expensive backend, usually a large language model call, and serves a stored response when an incoming query lands close enough in embedding space to one already answered. The promise: skip the call, return the same answer faster and for free. But “the same answer” is a claim about time as well as similarity, and most semantic caches never check the time part.

This matters most where a language model sits between a user and an answer about content that changes on its own schedule, for example monitoring how a brand is represented across AI-generated answers (answer engine optimization). The same style of question gets asked repeatedly, exactly the workload a semantic cache absorbs, but the facts behind it do not sit still. A cache that only asks “have I seen a similar question” and never “is that answer still true” keeps serving a stale answer indefinitely once written, with no way to tell a correct hit from a wrong one, because stale-hit-rate does not exist in standard cache evaluation.

Two gaps follow, motivating the two components I add. **No staleness concept in the serve path:** a cached answer about a fact with its own expiry schedule is served indefinitely, with “found in cache” silently treated as “correct.” **No cost term in eviction:** count-based policies such as LRU and LFU treat two entries with identical access patterns the same way even when one is far more expensive to regenerate, so a rarely accessed but costly entry can be evicted ahead of a cheap one merely accessed once more recently.

Industry and research reflect this gap unevenly. Frontier model providers do exact-prefix key-value caching with a wall-clock TTL and no similarity matching. Semantic-caching products use a single

fixed global threshold with plain TTL or count-based eviction, with no notion of per-topic staleness. The research literature is further ahead, proposing cost-aware and staleness-aware eviction formulas, but none names or measures stale-hit-rate; I compare against that literature, not the weaker product baseline.

This report makes three contributions. First, stale-hit-rate as a named, measured metric, and a semantic-index-based harness that exercises it under conditions closer to a real semantic cache than an oracle-cluster, exact-match one. Second, a diagnosed measurement failure rather than a negative result: cluster identity cannot be recovered from this workload’s query text by a general-purpose embedder, so per-cluster fitting was never actually exercised under real clustering, and I trace the exact cause (the query template differs across clusters only in two embedded numbers) rather than reporting the untestable comparison as if it settled the question. Third, a positive result on cost-aware eviction: given a fair test under conditions where eviction actually happens, which every earlier version of this suite failed to arrange, FreCoS’s eviction term measurably beats a frequency floor on cost saved. Section 2 justifies the baseline, Section 3 describes the design, Section 4 covers the workload and harness, and Section 5 reports seven experiments: one bracketing run and its two misspecification follow-ups, one five-row ablation, three parameter sweeps, and one gate-off cost-aware eviction test.

1.2 Related work

The closest prior art, and my primary eviction comparator, is Biton and Friedman’s work on semantic-cache eviction [1]. It proves optimal offline semantic-cache eviction is NP-hard, supplies polynomial-time heuristics, and presents online policies combining recency, frequency, and locality, evaluated on GPTCache with released code, part of why I chose GPTCache here. Two findings shape my design: LRU is weak on semantic workloads, and frequency-based policies are a stronger baseline, which is why every comparison in this report treats an LFU-class policy, not LRU, as the floor. FreCoS differs in adding two orthogonal signals, regeneration cost and content validity over time, and evaluating against a correctness metric, stale-hit-rate, their setting never defines.

The structurally closest published formula is Cortex’s LCFU policy [2], scoring an entry as log-scaled frequency times cost times latency times a staleness term, divided by size, with a TTL purge on top. The difference is narrow: Cortex’s staleness term is a static score assigned by a language model at write time, while FreCoS learns a decay rate per cluster from observed staleness; Section 5.1 tests whether that learning does measurable work. I did not reimplement Cortex’s formula, since it relies on remote tool-call latency metadata with no equivalent here.

Several other lines touch individual pieces without matching the combination: GDSF [3] ages access recency, not content validity; Category-Aware caching [4] assigns load-based, not learned, per-category TTLs (why I cut a third planned component, calibrated per-cluster thresholds, a lossy special case of this); FreshCache [5] puts staleness in the serve path as a probabilistic risk budget rather than a hard gate (future work here); SCALM [6] clusters entries similarly to my k-means step but with no staleness-aware significance score; MeanCache [7] is cost-motivated but does not target eviction; W-TinyLFU [8] is admission control, orthogonal to both my contributions. I considered and rejected vCache [9]: its contribution, per-prompt threshold learning with eviction explicitly out of focus, would have diluted rather than sharpened the comparison.

2 Baseline

2.1 Why GPTCache

I extended GPTCache [10] rather than a more actively maintained semantic cache, for four reasons. First, its most recent release (v0.1.44, August 2024) is frozen (maintainers state new integrations have stopped), which is a better scientific control than a moving target: the baseline measured early in this project is byte-identical to the one measured at the end. Second, architectural fit, verified by reading the vendored source directly: every component I touch sits behind a named interface, so my extension is purely additive (a metadata field via composition, a serving-path gate in one shared decision helper). FreCoS’s eviction policy is the one exception, added by monkeypatching `EvictionBase.get()` rather than through any extension hook, because that factory has none: a hardcoded if/elif chain with no lookup table a caller could add to. Verified end to end against a real `gptcache.manager.data_manager.SSDDataManager` in `tests/test_gptcache_integration.py`. I never edit vendored code. Third, GPTCache is the baseline the closest prior art already uses, letting me in principle run Biton and Friedman’s released policy

directly, though in practice I substituted a documented proxy (Section 5.2). Fourth, its default embedding backend runs CPU-deterministically, keeping every run reproducible without a GPU. One trade made deliberately: GPTCache’s README states new contributions are not merged, foreclosing this course’s upstream-adoption bonus.

2.2 GPTCache’s behavior and default eviction policy

GPTCache serves a cached response when a query is semantically close enough to one already seen, using a single global cosine threshold (0.8 by default). Its eviction policies are purely count-based (LRU, LFU, FIFO, random); none reads generation time, regeneration cost, or retains metadata on an evicted entry, though creation and last-access timestamps are persisted with no downstream consumer, precisely the gap my extension repurposes. There is no admission control, and no cost, false-hit, or staleness tracking anywhere in the built-in reporting layer.

3 Design

3.1 Contribution

The contribution is one learned artifact with two consumers. FreCoS fits a per-cluster staleness model and uses it at both ends of the cache’s lifecycle: as a hard validity gate when serving a candidate hit, and as a soft decay term when choosing an eviction victim. The metric that makes both uses visible is stale-hit-rate: the fraction of served hits whose content was already invalid when served.

Cluster assignment comes from a fixed, seeded k-means over cached-query embeddings from GPTCache’s default ONNX model, fit once per trace on the calibration split, with every row assigned to its nearest centroid. This is a change from an earlier version, which read cluster identity directly off the workload generator’s ground truth; every number in every earlier version of this suite was produced under that oracle-perfect assignment. Section 5.1 reports what changes once clustering has to work from text, the report’s central finding. The generator’s true label is still recorded on every row, used only to compute the adjusted Rand index against the learned assignment (`cluster_ari`, reported alongside every bracket-style result), never consumed by the gate or eviction. Both consumers use the same (possibly wrong) cluster identity and neither recomputes it independently; I checked directly for disagreement between them in the ablation (Section 5.2) and found none.

For each cluster, a decay rate λ_c is fit from observed staleness in a held-out calibration split, assuming exponential validity decay: $\Pr[\text{still valid after } \Delta t] = e^{-\lambda_c \Delta t}$. A TTL is derived at confidence level c as $\text{ttl} = -\ln(c)/\lambda_c$. Three modes produce this table with the same output shape: *learned* fits each cluster independently by maximum likelihood, falling back to a pooled rate below thirty observations; *global* pools every cluster regardless of sample size; *oracle* takes the rate from the generator’s ground truth, an upper bound never available in a real deployment. Table 1 lists every tunable parameter, its default, and where its effect is measured.

The validity gate is the first consumer: on a candidate hit, if the entry’s age exceeds its cluster’s TTL, it is treated as a miss and the event counts as a stale hit prevented rather than served. This checks creation time only, never last-access time: staleness is a property of when an answer was generated, not how recently it was asked for, so a popular but outdated entry is not treated as fresh just because it is popular.

The eviction value function is the second consumer:

$$\text{value} = \log(2 + \text{freq}) \cdot \log(1 + \kappa \cdot \text{regen_cost}) \cdot e^{-\lambda_c \cdot \text{age}} \quad (1)$$

where age is from creation time, never last access; $\kappa = 1000$ log-compresses regeneration cost so its heavy tail cannot dominate; the victim is the lowest-value entry, ties broken by oldest creation then lowest entry id. An earlier version also divided by `size_bytes` to evict larger entries first; that term is removed, since eviction here runs under an entry-count budget, not a byte budget: evicting a 10-byte and a 10,000-byte entry both free exactly one slot, so size-normalization has no budget-economics meaning. This was confirmed empirically at the highest-eviction-pressure cache-size point in the project’s sweep (`results/ablation/size_term_isolation/`): FreCoS with and without the size term was statistically indistinguishable on every metric. A byte-budgeted eviction loop, where size would have real meaning, remains future work.

One detail came from a design mistake a property test caught, not one found by inspection: my original frequency term, $\log(1 + \text{freq})$, is exactly zero for a never-accessed entry, zeroing the whole product

regardless of cost or freshness, so a freshly written entry would always be evicted next. Appendix A’s property test failed against the literal formula and caught this; I fixed it to $\log(2 + \text{freq})$, same concave shape but strictly positive at zero frequency, applied throughout.

Because the gate already refreshes anything past its TTL, the decay term is not doing staleness prevention on its own: it is a soft prior among candidates the gate has already judged fresh. The ablation (Section 5.2) tests this directly.

Table 1: Every tunable parameter, its default, and where its effect is measured in Section 5. **Config** ships $c = 0.95$; the bracketing and ablation experiments in Sections 5.1–5.2 run at $c = 0.90$, called out wherever those results appear.

Parameter	Default	Swept	Measured in
lambda source (mode)	learned	global, learned, oracle	5.1
TTL confidence c	0.95	0.80, 0.90, 0.95, 0.99	5.4
cluster count K	10	5, 10, 20, 50	5.4
cache size (entries)	1650 (see note)	5%–80% of working set	5.4
κ (cost log-compression)	1000	not swept	—
calibration fallback threshold	30 obs.	not swept	—

Cache size is 1650 entries in the bracketing and ablation experiments; the cache-size sweep itself spans 5%–80% of the working set and identifies 990 entries as the knee (Section 5.4). κ and the thirty-observation fallback threshold were fixed at values reasonable at design time and not independently swept.

4 Experimental setup

4.1 Workload

All experiments run on W1, a deterministic, seeded synthetic workload generator built for this project. Before other implementation work, I timeboxed a feasibility spike into a second workload derived from Wikipedia revision history: an automatic sentence-change rule agreed with hand labels only 40% of the time against an 80% go/no-go bar set in advance (full spike in `docs/w2-feasibility.md`), so I proceeded with W1 alone.

W1 generates queries across several streams inside each topic cluster: canonical queries establishing a ground-truth answer, near-duplicate paraphrases, novel long-tail queries that never repeat, and time-shifted repeats. Every cluster draws its own half-life and regeneration-cost distribution, and at least one repeat per answer lands past its true expiry, giving the gate something to reject. Each trace splits by arrival time, 30% calibration and 70% evaluation, so no parameter is fit on data it is later scored against.

The default generator draws half-life from an exponential distribution, matching the staleness fitter’s own assumption exactly. Section 5.1.1 reruns the bracketing experiment with that assumption deliberately violated: once with half-life drawn from a Weibull distribution (same mean, different shape), and once, the stronger check, from a mixture of two exponentials within each cluster.

The generator’s two ground-truth distributions, regeneration cost and content half-life, were meant to be calibrated against the Azure LLM Inference Trace and observed update-frequency in a public corpus. Neither was reachable from my development environment, so both are parameterized from commonly cited public figures instead: a plausibility check, not a real external fit, and a limit on external validity.

4.2 Simulation model

The benchmark harness is a pure trace replayer; no language model is ever called. On a miss, response content and cost are whatever the trace recorded, and miss latency is simulated from a seeded log-normal scaled by entry size, deterministic per query: a documented placeholder, not fit to a real trace. Hit latency and extension overhead (index lookup plus gate check) are measured for real: median overhead is 0.47ms per decision (Table 5), roughly three orders of magnitude higher than an earlier exact-match-index version, since a real cosine search over the cache’s embedding matrix costs far more than a dict lookup. Reported latency is a hybrid of real and simulated cost, the right trade for keeping runs free and deterministic, but a limitation worth stating.

Two limitations that shaped every earlier version’s absolute numbers are gone as of this version: the lookup index is now `benchmarks.semantic_index.SemanticIndex`, brute-force cosine similarity at GPTCache’s default 0.8 threshold, not exact-text-match; and cluster identity comes from real k-means over query embeddings, not the generator’s ground truth (Section 3.1). Neither is free of new limitations: a real semantic index makes false-hit-rate a live metric for the first time, and it is stark: this workload’s canonical query template differs across clusters only in two embedded numbers, which a real embedder scores at 0.6–0.9 cosine similarity regardless of true identity, above the 0.8 threshold, so roughly 97% of served hits across every bracket-style experiment are false. This is a property of this synthetic workload’s text, not the cache mechanism, but every hit-rate number below should be read alongside its false-hit-rate (Section 5.1 carries both). Real clustering similarly changes every stale-hit-rate number relative to any earlier version of this report.

The cache starts empty in every run. The first 10% of the evaluation split warms the cache and is excluded from every metric.

4.3 Metrics and statistics

Stale-hit-rate is the fraction of served hits where serve time exceeds the entry’s true expiry; false-hit-rate is the fraction where the served answer’s identity does not match the query’s; useful-fraction-of-hits is the fraction of served hits, not all scored queries, that are both correct and non-stale: $(n_{\text{hits}} - n_{\text{stale}} - n_{\text{false}})/n_{\text{hits}}$ (an earlier, all-queries definition breaks once false-hit-rate is large, since it can subtract more than n_{hits} and go negative; per-hit is also the more honest question: given a hit happened, was it worth anything). Cost saved sums regeneration cost over non-stale hits only, and does not also exclude false hits, so at this workload’s false-hit-rate (roughly 92–97%, Section 4) the reported cost-saved numbers are dominated by hits that returned the wrong answer; `cost_saved_usd` should be read as an upper bound, not money actually saved by correct answers, until the metric excludes false hits too and every affected experiment is rerun. Adjusted Rand index (`cluster_ari`), reported alongside every bracket-style result, measures how well the learned assignment matches the true cluster labels: 1.0 is a perfect match, 0.0 is what a random assignment produces in expectation.

Two primary comparisons are designated up front and reported without correction: learned versus global (does per-cluster learning beat pooling), and gate-on versus the LFU floor (does the gate do measurable work). Every other Mann-Whitney test is secondary, corrected via Holm-Bonferroni across the secondary family, raw and adjusted p -values reported together (`analysis/multiple_comparisons.py`, hand-rolled for the same reason Mann-Whitney U is).

Every configuration ran ten independent seeds, each with its own freshly generated trace. Results are medians with 95% CIs from a percentile bootstrap over the ten per-seed values (10,000 resamples, fixed seed), with raw hit and stale-hit counts alongside every rate since a ten-seed bootstrap CI is coarser than it looks. Group comparisons use Mann-Whitney U with rank-biserial correlation as effect size, since at $n = 10$ the smallest achievable two-sided p is about 0.0002 and cannot alone distinguish a large effect from a small one. A non-significant result is an absence of detected difference, never equivalence; where I claim two conditions behave alike, raw counts back it up, not just a large p -value.

Every experiment ran on the machine recorded in its own `env.json`: an Apple M2 Pro (10 physical/10 logical cores), 16 GB RAM, macOS (Darwin 25.5.0), Python 3.14.2. This was a shared development machine, not a dedicated benchmarking host, which matters for the peak-RSS finding this report already flags as unreliable (Table 5).

5 Results

5.1 Bracketing: does learned staleness beat naive pooling

This experiment tests the core claim behind fitting a staleness rate per cluster at all: does it beat a rate pooled across every cluster, and how far is it from the true rate a real system could never observe. I ran three lambda sources, global, learned, and oracle, for 10 seeds each on the same 12,000-query W1 configuration, gate enabled at $c = 0.90$, FreCoS eviction throughout. Cluster identity comes from real k-means over query embeddings (Section 3.1), not the generator’s true label; the oracle arm alone still uses the true label internally, since its lambda table is keyed by true cluster identity and has no other meaning once clustering is imperfect.

Table 2: Under real clustering, learned no longer tracks oracle: it tracks global. The primary comparison (learned vs. global) is not significant; the secondary comparison against oracle is, with a large effect, in the wrong direction for the per-cluster fitting claim this experiment was designed to test. `cluster_ari` is 0.036 across every row, barely above what a random assignment produces. Median `n_hits` out of 7560 scored queries is 4333.5/4252.0/4294.5 for global/learned/oracle; Mann-Whitney vs. learned is $U = 49.0$, $p \approx 0.940$ (primary, n.s.) for global and $U = 96.0$, $p \approx 0.0005$, adj. $p \approx 0.0046$, $r = 0.78$ for oracle.

Lambda source	Median stale-hit-rate	95% CI	Median cost saved (USD)	cluster_ari	False-hit-rate
Global	0.077	(0.057, 0.097)	14.86	0.036	0.969
Learned	0.071	(0.059, 0.091)	14.95	0.036	0.969
Oracle	0.042	(0.039, 0.047)	15.51	0.036	0.968

Learned and global, the primary comparison, are statistically indistinguishable ($U = 49.0$, $p \approx 0.940$, $r = 0.02$): pooling every cluster’s observations into a single rate performs the same as fitting each cluster’s own. Learned and oracle, a secondary comparison, separate significantly with a large effect in the direction that per-cluster fitting is *worse* than an oracle rate keyed to true cluster identity ($U = 96.0$, $p \approx 0.0005$, adjusted $p \approx 0.0046$, $r = 0.78$), the opposite of an earlier, oracle-cluster version, where learned tracked oracle almost exactly and beat global.

The reason is clustering quality, not calibration or the maximum-likelihood fit. Median `cluster_ari` across all 30 runs is 0.036, close to the 0.0 a random assignment produces and far from a perfect match. Direct inspection explains why: W1’s canonical query template (“What is the current status of topic {cluster}-{answer}?”) is identical across every cluster except two embedded numbers. Ten canonical queries drawn from ten different clusters, embedded with the same ONNX model, score 0.6–0.9 cosine similarity against each other, often higher than genuine same-cluster paraphrase pairs. A general-purpose embedder is not sensitive to numeric substitution inside an otherwise fixed template the way it would need to be for cluster identity to be recoverable at all. This is a property of how W1’s text was generated, not a bug in the k-means implementation or embedder wiring: k-means converges normally on the embeddings it is given, and it is the embeddings themselves that carry almost no cluster signal.

Three follow-ups (a scarcer-calibration bracket and two misspecification checks, Section 5.1.1) all reproduce the identical pattern: learned tracks global, not oracle, with `cluster_ari` in the same 0.036–0.062 range. This rules out calibration sample size and half-life-distribution shape as confounds; the clustering step, which happens before either in the pipeline, is where the per-cluster advantage is lost. Figure 1 plots stale-hit-rate for all three lambda sources across the well-specified, calibration-sparse, and adversarial-mixture conditions.

5.1.1 Misspecification: what happens when the exponential assumption is wrong

The first check drew W1’s true half-life from a Weibull distribution instead of exponential (same mean, shape 2); the second, stronger check drew it from a 50/50 mixture of two exponentials within each cluster, biased against the fitter’s single-rate maximum-likelihood estimate. Both were run to test whether the fitter’s exponential assumption mattered, independent of clustering; under real clustering, both instead reproduce the well-specified bracket’s learned-tracks-global pattern.

Table 3: Both misspecification checks reproduce the well-specified bracket’s pattern: learned tracks global (not significant), not oracle (significant, large effect), with effect sizes close to the well-specified run. Weibull: learned vs. global $U = 47.0$, $p \approx 0.82$ (n.s.); learned vs. oracle $U = 100.0$, $p \approx 0.0002$, adj. $p \approx 0.0023$, $r = 0.85$. Mixture: learned vs. global $U = 44.0$, $p \approx 0.65$ (n.s.); learned vs. oracle $U = 99.0$, $p \approx 0.0002$, adj. $p \approx 0.0026$, $r = 0.83$. `cluster_ari` is 0.04–0.06 for both, matching the well-specified run.

Check	Lambda source	Median stale-hit-rate	Median cost saved (USD)	cluster_ari
Weibull	Global	0.027	15.54	0.041
Weibull	Learned	0.028	15.57	0.041
Weibull	Oracle	0.006	16.01	0.041
Mixture	Global	0.151	14.20	0.037
Mixture	Learned	0.142	14.43	0.037
Mixture	Oracle	0.095	14.92	0.037

Under real clustering, the half-life distribution’s shape is a second-order effect: the dominant factor

is clustering quality, untouched by how `half_life` is drawn, since cluster identity is determined entirely by the query-text embedding step. This rules out calibration sample size and model misspecification as confounds for the brackets finding, confirming clustering as the bottleneck regardless of the ground-truth half-life distribution.




Figure 1: Stale-hit-rate under global, learned, and oracle lambda sources, at the original calibration scale and a roughly sevenfold sparser scale, both under real clustering. Learned tracks global at both scales, not oracle, reversing the oracle-cluster-era version of this figure.

5.2 Ablation: isolating the gate and the eviction term

I ran five configurations for 10 seeds each at TTL confidence $c = 0.90$: stock LRU and stock LFU gate-disabled (LFU is the floor, given Biton and Friedman’s finding that LRU underperforms on semantic workloads), a documented substitute for Biton and Friedman’s released policy also gate-disabled, and the validity gate combined with plain LFU or full FreCoS eviction. A sixth row from an earlier version, FreCoS with the size-normalization term removed, no longer exists: that term was removed from Equation 1 entirely (Section 3.1).

Table 4: Median stale-hit-rate, hit rate, and false-hit-rate by row, 10 seeds each. Row 2 (LFU, gate off) is the floor. Rows 1–3 are byte-identical seed by seed; rows 4–5 are byte-identical seed by seed, both traced to real causes below.

Row	Configuration	Stale-hit-rate	Hit rate	False-hit-rate
1	LRU, gate off	0.720	0.952	0.961
2	LFU, gate off (floor)	0.720	0.952	0.961
3	B&F substitute, gate off	0.720	0.952	0.961
4	LFU, gate on	0.071	0.562	0.969
5	FreCoS, gate on (full stack)	0.071	0.562	0.969

The gate’s own effect (the second primary comparison, row 4 against the floor) is large and significant: $U = 0.0$, $p \approx 0.00016$, $r \approx 0.85$, cutting stale-hit-rate roughly 10-fold. That part of the original finding survives the rerun to real clustering and a semantic index intact. The gate is a per-entry age check against a per-cluster TTL, and works correctly whether or not the cluster feeding it is right, because the same assignment is used consistently at calibration and serve time.

Rows 1–3 tie because, at this configuration’s real semantic hit rate (95%), only 339–372 misses occur out of 7560 scored queries per run, under the 1650-entry budget, so eviction never runs. LRU, LFU, and the B&F substitute all tie because none is ever called to select a victim, not because they agree on one. Section 5.3 confirms this directly by rerunning this exact gate-off configuration at a cache size small enough to force real eviction, finding real, significant differences between the three policies there. Rows 4–5 tie for an unrelated reason: FreCoS’s decay term contributes nothing measurable once the gate is active, discussed below.

Isolating the gate versus the decay term: gate alone (row 4 vs. row 2) gives $\Delta = -0.649$ ($U = 0.0$, $p \approx 0.00016$, $r \approx 0.85$); FreCoS’s decay term on top of an already-active gate (row 5 vs. row 4) gives $\Delta = 0.000$, byte-identical. The gate accounts for the entire measurable stale-hit-rate improvement here, matching the design’s own framing of the decay term as a soft prior: with the gate active, every entry eviction sees is already known-fresh, so the decay term has little left to differentiate on.

Table 5: Latency and resource use at the LFU floor (row 2) and the full stack (row 5). Extension overhead is roughly three orders of magnitude higher than an earlier, exact-match-index version, since the semantic index performs a real cosine search. Latency/throughput are dominated by the simulated miss-cost distribution (Section 4.2), a single run with no CI, unlike every other table here.

Metric	Row 2 (floor)	Row 5 (full stack)
Extension overhead, mean (ms), measured	0.468	0.478
Latency, mean (ms), simulation-dominated	7.01	58.93
Latency, p95 (ms), simulation-dominated	1.11	225.37
Latency, p99 (ms), simulation-dominated	182.86	366.09
Throughput (queries/s), simulation-dominated	142.7	17.0
Peak RSS (MB), median (95% CI)	650.6 (632.8, 716.3)	649.5 (631.7, 688.2)
CPU (%), measured	94.3	94.4

Peak RSS is statistically indistinguishable between rows, unsurprising since both run through the same real index on a shared development machine, not a dedicated benchmarking host. The mean-latency and p95/p99 gap (7ms to 59ms mean) is large because the gate converts a large fraction of hits into misses at this real hit rate, and a simulated miss costs far more than a real cache-side decision. Still simulation-dominated, not a claim about real end-to-end latency.

No latency CDF accompanies Table 5, though the guideline asks for one: miss latency is drawn from a seeded log-normal placeholder (Section 4.2), not measured from a real backend, so its distribution plot would show the shape of a random-number generator, not any real system’s latency. The one measured, not simulated, latency quantity (extension overhead) does not vary enough across ten seeds to make a distribution plot informative beyond the mean already reported.




Figure 2: Stale-hit-rate across all five ablation rows, with 95% confidence intervals. Row two, LFU with the gate disabled, is the floor of comparison.

5.3 A fair test of cost-aware eviction

Every ablation above ran with the gate on, which pre-filters stale entries before eviction ever sees them, leaving FreCoS’s cost and decay terms nothing to differentiate on. This experiment runs gate off (Null-Gate), three eviction policies (FreCoS, LFU, LRU) x 10 seeds, at a cache size (100 entries) chosen after measuring that the main ablation’s 1650-entry cache never fills under this workload’s real hit rate (only 339–474 distinct entries are ever needed with an unlimited cache). At 100 entries, 1300+ misses occur per run, forcing real eviction pressure.

Table 6: With real eviction pressure, FreCoS beats the LFU floor on cost saved with a large, significant effect: the first result here where FreCoS’s cost-awareness shows measurable work. LRU strictly dominates FreCoS: more cost saved and a lower stale-hit-rate, with no per-seed overlap on either metric.

Policy	Median cost saved (USD)	95% CI	Median stale-hit-rate
FreCoS	11.83	(9.17, 13.46)	0.491
LFU	7.17	(5.52, 8.27)	0.676
LRU	17.09	(14.30, 21.06)	0.227

FreCoS vs. LFU on cost saved: $U = 94.0$, $p \approx 0.0009$, $r \approx 0.74$ (significant, FreCoS higher). FreCoS vs. LRU on cost saved: $U = 11.0$, $p \approx 0.0032$, $r \approx 0.66$ (significant, LRU higher). On stale-hit-rate, LRU is lower with no per-seed overlap (LRU’s worst seed, 0.326, still beats FreCoS’s best, 0.460): LRU wins on both metrics. This is not because LRU has any staleness mechanism (with no gate active it has none), but because on this trace, recency of access correlates with recency of write, so evicting the least-recently-accessed entry incidentally tends to evict the entry most likely to already be stale. LRU gets a freshness benefit for free, no staleness model, no calibration split, no gate, which is a genuine challenge to this project’s premise that cost- and staleness-awareness are needed at all. What does hold up: FreCoS beats the pure-frequency LFU floor on cost saved with a large effect, the first experiment where FreCoS’s cost term shows real, significant work. It does not establish that FreCoS beats LRU, and the report does not claim otherwise.

5.4 Cache-size sweep and the correctness-efficiency trade-off

I swept cache size across 5 points spanning 5%–80% of the workload’s distinct-answer working set, full stack throughout. Hit rate reaches its ceiling by 15% of the working set (990 entries) and flattens completely from there; every point from 990 entries on produces identical per-seed hit counts. The knee falls at 990 entries, smaller than an earlier exact-match-index sweep found (1980 entries), since the semantic index’s higher baseline hit rate means fewer distinct entries need to be held to saturate it. This cache size anchors every other sweep in this report.

The next sweep holds cache size fixed at that knee and sweeps the gate’s target confidence level.

Table 7: Useful-fraction-of-hits (correct and non-stale, among served hits only; see Section 4.3) is zero or indistinguishable from zero through confidence 0.90, small but non-zero at 0.95, and only substantial at 0.99, moving in lockstep with false-hit-rate. Stale-hit-rate falls 18-fold across the same range.

TTL conf.	Med. stale-hit-rate	Med. hit rate	Med. useful-frac.-of-hits	Med. false-hit-rate
0.80	0.133	0.684	0.000	0.975
0.90	0.071	0.562	0.000	0.969
0.95	0.039	0.473	0.000 (CI up to 0.008)	0.962
0.99	0.007	0.269	0.068	0.920

Useful-fraction-of-hits stays at or indistinguishable from zero through 0.90, still tiny at 0.95, and reaches only 0.068 at 0.99 (0.95 vs. 0.99: $U = 0.0$, $p \approx 0.00013$, adjusted $p \approx 0.0023$, $r \approx 0.86$, significant). This is a direct consequence of false-hit-rate sitting above 0.92 at every confidence level: almost every served hit is false regardless of TTL strictness, since false hits come from the semantic index’s threshold match, not the gate. Raising confidence still buys a real, large stale-hit-rate reduction (18-fold), but “the gate is nearly free,” an earlier version’s central claim, does not survive a real semantic index: this workload’s false-hit problem dominates correctness at every confidence level.

A third sweep varied cluster count at 5, 10, 20, and 50. `hit_rate` and `stale_hit_rate` remain non-monotone with heavily overlapping CIs, matching an earlier exact-match-era version. `cluster_ari` itself rises steadily with K (0.022 at $K = 5$ to 0.055 at $K = 50$): more, smaller ground-truth clusters are marginally easier for k-means-on-embeddings to separate, though every value tested stays close to what a random assignment produces. This result sits in a supplementary table (`analysis/figures/supplementary.csv`) rather than a headline figure, since it clears neither the effect-size nor confidence-interval bar the rest of this report’s figures were held to.




Figure 3: Hit rate against cache size, with the identified knee at 15% of the working set (990 entries) marked.




Figure 4: Stale-hit-rate and useful-fraction-of-hits against the gate’s target confidence level. Stale-hit-rate falls 18x; useful-fraction-of-hits stays near zero through 0.95 and only becomes substantial at 0.99, tracking false-hit-rate’s decline.

6 Discussion

The central finding, Section 5.1, is a diagnosed measurement failure, not a verdict on per-cluster learning: there were no real clusters for it to exploit, so the ARI numbers there do not license a general claim. An earlier version reported the opposite conclusion under oracle-perfect cluster identity; that framing was not wrong given what it tested, but this remediation shows the question (does per-cluster fitting recover a rate well, given correct clusters) is not yet answerable end to end, because the harness meant to supply real clusters cannot do so on this workload’s text.

What survives: the validity gate still works correctly whether or not the cluster feeding it is accurate, since the same cluster is used consistently at calibration and serve time (Section 5.2). And FreCoS’s cost-aware eviction term, tested here for the first time under conditions where eviction actually happens, beats a frequency floor with a large effect (Section 5.3), a genuinely positive result no earlier version of this suite tested for. It does not beat plain LRU, which strictly dominates FreCoS in the same test, because recency of access happens to correlate with recency of write on this trace, so LRU gets a freshness benefit for free with no staleness model at all, a direct challenge to the premise that cost- and staleness-awareness are needed on recency-correlated workloads.

A related finding shares the same root cause: this workload’s construction cannot support a meaningful correctness evaluation once a real semantic index is used, because the canonical query template differs across clusters and answers only in two embedded numbers, which a real embedder cannot use to distinguish them. Roughly 92–97% of hits served across every experiment, at every TTL confidence tested, are false hits: the cache confidently returns a real cached answer to a different question than the one asked. This is a claim about this workload’s text, designed for an oracle-cluster world that never needed the text to be discriminable, not about semantic caches generally. Every number in Section 5 should be read next to its false-hit-rate until the text is redesigned or a real corpus replaces it.

A further, structural limitation present regardless of clustering or the index: stale-hit-rate is scored against `valid_until`, the same ground-truth field that supplies the oracle arm’s lambda in Section 5.1. The metric, the ground truth it is checked against, and the ceiling one arm of every bracket comparison is measured relative to, all come from one self-authored generator. This does not invalidate relative comparisons within a single experiment, since both sides are scored against the same ground truth equally, but no number here should be read as a claim about staleness detection against an independent, external notion of correctness. I also attempted external ground truth from Wikipedia and stopped after a feasibility spike measured 40% rule agreement against an 80% bar set in advance. Neither prior-art comparator, Biton and Friedman’s released policy nor Cortex’s LCFU, was run against this workload, so the eviction axis is measured only against count-based policies and this project’s own FreCoS variant.

7 Conclusion and future work

I set out to name a gap in how semantic caches are evaluated, stale-hit-rate, and build a system that closes it, with two components matching the two gaps in Section 1.1. Neither component’s story is the one I expected: the gate holds up regardless of clustering quality, the per-cluster learning meant to sharpen it was never actually put to the test its own harness intended (Section 6), and the eviction term’s cost-awareness, previously reported as contributing nothing, turns out to do real work once tested under conditions where eviction actually happens, though plain LRU still beats it here for an unmodeled reason (Section 5.3).

The most consequential next step is redesigning W1’s query text so cluster and answer identity are recoverable from a real embedder, not just a generator’s internal label: the current template’s reliance on embedded numbers alone breaks both clustering and the semantic index, from the same root cause. This fix is fully specified (a distinct topic phrase per cluster, a distinct entity phrase per answer_id, replacing the two-number template in `workloads/w1_synthetic/generator.py`’s canonical-text line) but not yet applied or rerun: the Phase 2 rerun that produced this version’s numbers already took just under 9 hours, and changing the generator would require rerunning every experiment against new committed CSVs, which is future work, not a result claimed here. The same applies to excluding false hits from `cost_saved_usd` (Section 4): the metric change is prepared, not yet rerun. Running Biton and Friedman’s policy and Cortex’s LCFU directly would close the comparator gap now that a real index exists to run them against. Calibrating the generator against a real inference trace would strengthen external validity, and breaking the circularity between stale-hit-rate and its generator-supplied ground truth would be the strongest single improvement here. Dynamic re-clustering, a FreshCache-style probabilistic staleness budget, byte-budgeted eviction (where the size term removed from Equation 1 would have real meaning), and admission control remain future work.

Appendix A: correctness evidence

Correctness is checked at four levels. A reference oracle (`tests/oracle/reference_cache.py`) implements the same hit/miss/stale decision as the extension, using a naive dict-and-linear-scan approach with no shared code, and every pipeline decision is cross-checked against it in `tests/test_stock_parity.py`. A callable invariant suite (`tests/invariants.py`) runs inside both unit tests and `benchmarks/harness.py` on every experimental run, checking budget bounds, argmin eviction selection, that no entry is served past its cluster’s TTL, and that staleness decisions derive from creation time, never last-access time. Property-based tests in `tests/test_freco.py` check that the eviction value function is monotonic in each remaining term independently. A fourth level, `tests/test_gptcache_integration.py`, drives a real `gptcache.manager.data_manager.SSDDataManager` through enough writes to force eviction and confirms the victim GPTCache’s own eviction factory selects matches `FreCoSEviction.select_victim()` computed directly from the same metadata. This confirms FreCoS is reachable through GPTCache’s real eviction path, not merely that the scoring function is correct in isolation.

The most persuasive of these artifacts is a property test that failed against my original design rather than a bug found by inspection: `test_cold_start_entry_scores_above_zero` in `tests/test_freco.py` asserts a fresh entry must score above zero, and failed against the literal $\log(1 + \text{freq})$ formula: how I found and fixed the $\log(1 + \text{freq}) \rightarrow \log(2 + \text{freq})$ issue in Section 3.1.

Appendix B: code and data artifacts

This repository is public at <https://github.com/LiorLotan2/freco>. A `Dockerfile` and `environment.yml` (or `requirements.txt` for a plain venv) build a working environment from a clean clone, and `.github/workflows/ci.yml` reruns the full test suite, a benchmark smoke test, an experiment-path smoke test, and a figure-consistency check on every commit. `benchmarks/samples/smoke_row.csv` and `smoke_row.json` are committed sample benchmark logs reproducible with `make bench-smoke`. `make experiments` reruns all seven experiments behind Section 5 from a clean clone in dependency order; `make figures` regenerates every figure in this report from the resulting CSVs.

Table 8: Where every artifact referenced in this report lives in the repository.

Artifact	Location
Pipeline seam (serve-path decision)	<code>gptcache_ext/pipeline.py</code>
Additive metadata storage	<code>gptcache_ext/metadata.py</code>
Staleness model (fitter, gate)	<code>gptcache_ext/staleness/</code>
Real clustering (embedder, k-means, ARI)	<code>gptcache_ext/staleness/embedder.py</code> , <code>clusters.py</code> , <code>assign_real_clusters.py</code> , <code>cluster_accuracy.py</code>
Eviction (FreCoS and baselines)	<code>gptcache_ext/eviction/</code>
Semantic index (cosine, 0.8 threshold)	<code>benchmarks/semantic_index.py</code>
Synthetic workload generator (W1)	<code>workloads/w1_synthetic/</code>
Benchmark harness and metrics	<code>benchmarks/harness.py</code> , <code>metrics.py</code>
Experiment runners (Section 5)	<code>benchmarks/runners/</code>
Raw results, summary, and environment per experiment	<code>results/*/results.csv</code> , <code>summary.md</code> , <code>env.json</code>
Reproduce all experiments (one command)	<code>make experiments</code>
Figure regeneration (one command)	<code>make figures</code>
Multiple-comparison correction	<code>analysis/multiple_comparisons.py</code>
Reference oracle	<code>tests/oracle/reference_cache.py</code>
Invariant suite	<code>tests/invariants.py</code>
Property tests	<code>tests/test_freco.py</code>
Real GPTCache eviction-factory integration test	<code>tests/test_gptcache_integration.py</code>
GPTCache pinned commit and source-map	<code>docs/baseline-source-map.md</code>
Wikipedia feasibility spike	<code>docs/w2-feasibility.md</code>
Docker / environment for a clean rebuild	<code>Dockerfile</code> , <code>environment.yml</code> , <code>requirements.txt</code>
Continuous integration	<code>.github/workflows/ci.yml</code>
Sample benchmark logs	<code>benchmarks/samples/</code>
Per-phase change log	<code>CHANGES.md</code>

No figure in this report is hand-edited; every one is regenerated directly from a committed results CSV by `make figures` (Table 8).

References

- [1] D. D. Biton and R. Friedman. *From Exact Hits to Close Enough: Semantic Caching for LLM Embeddings*. arXiv:2603.03301, 2026.
- [2] C. Ruan, C. Bi, K. Zheng, Z. Shi, X. Wan, and J. Li. *Cortex: Achieving Low-Latency, Cost-Efficient Remote Data Access for LLM via Semantic-Aware Knowledge Caching*. arXiv:2509.17360, 2025.
- [3] L. Cherkasova and G. Ciardo. *Role of Aging, Frequency, and Size in Web Cache Replacement Policies*. Proceedings of the International Conference on High-Performance Computing and Networking (HPCN), 2001.
- [4] C. Wang, X. Liu, Y. Zhu, A. Youssef, P. Nagpurkar, and H. Chen. *Category-Aware Semantic Caching for Heterogeneous LLM Workloads*. arXiv:2510.26835, 2025.
- [5] M. Mansoor, T. Ahmad, and Y.-C. Yoon. *Risk-Constrained Freshness-Aware Semantic Caching for Open-Web Retrieval-Augmented LLMs*. arXiv:2607.04281, 2026.
- [6] J. Li, C. Xu, F. Wang, I. M. von Riedemann, C. Zhang, and J. Liu. *SCALM: Towards Semantic Caching for Automated Chat Services with Large Language Models*. arXiv:2406.00025, 2024.
- [7] W. Gill, M. Elidrisi, P. Kalapatapu, A. Ahmed, A. Anwar, and M. A. Gulzar. *MeanCache: User-Centric Semantic Caching for LLM Web Services*. arXiv:2403.02694, 2024.
- [8] G. Einziger, R. Friedman, and B. Manes. *TinyLFU: A Highly Efficient Cache Admission Policy*. arXiv:1512.00727, 2015.
- [9] L. G. Schroeder, A. Desai, A. Cuadron, K. Chu, S. Liu, M. Zhao, S. Krusche, A. Kemper, M. Zaharia, and J. E. Gonzalez. *vCache: Verified Semantic Prompt Caching*. arXiv:2502.03771, 2025. ICLR 2026.
- [10] Zilliz. *GPTCache: A Library for Creating Semantic Cache for LLM Queries*. <https://github.com/zilliztech/GPTCache>, 2024. Version 0.1.44, commit bae7ffeef774e762d9d4e60fce70be00011188a6.